

Objective:

The task is to develop a small web service that integrates with OpenAI's GPT-4 API and exposes an endpoint for querying the GPT model. The service should process user input, modify the prompt dynamically based on user type, and return the response in a structured format.

Task Overview:

1. Develop a Backend Service (using Python and FastAPI):
 - Create a REST API with at least one POST endpoint that accepts a JSON payload with a user query.
 - The API should interact with OpenAI's GPT-4 API to generate responses.
2. Role-Based Behavior:
 - The user should provide their role (e.g., "admin" or "user") in the request.
 - Based on the role, modify the prompt to change how the model responds. For example:
 - Admin: The GPT response should be more formal and provide more detailed explanations.
 - User: The GPT response should be simpler and more casual.
3. Query Parameters:
 - The request payload should include:
 - role: a string to define the user role (admin/user).
 - query: a string representing the user's question or input.
4. Response Structure:
 - The API response should be structured as follows:

```
{
  "original_query": "User's original query",
  "gpt_response": "GPT-generated response",
  "timestamp": "Time when the response was generated"
}
```

5. Bonus:
 - Implement logging that tracks the original query, GPT's response, and the user role.
 - Add error handling for potential issues such as missing inputs, invalid API keys, or network errors.

Instructions for Submission:

- Language: Use Python for the backend.
- Framework: FastAPI for building the API.
- OpenAI API: Integrate with the GPT-4 API using OpenAI's official library.
- Documentation: Provide clear instructions for how to:
 - Set up the environment.
 - Run the application locally.
 - Test the API using sample curl commands or a tool like Postman.
- Time Estimate: ~2-4 hours.
- Bonus Tasks: Optional but highly recommended to attempt for extra credit.

Sample Input and Output:

Request:

```
{
  "role": "admin",
  "query": "What is the current status of artificial intelligence?"
}
```

Response:

```
{
  "original_query": "What is the current status of artificial intelligence?",
  "gpt_response": "As of 2024, artificial intelligence continues to advance in various fields, including natural language processing, robotics, and healthcare...",
  "timestamp": "2024-10-07T10:15:30Z"
}
```

Evaluation Criteria:

1. API Functionality (40%):
 - Does the API work as expected? Is it able to correctly handle role-based modifications to the prompt?
 - Are the responses from GPT accurate and well-structured?
2. Code Quality (25%):
 - Is the code clean, modular, and well-organized?
 - Is there proper error handling and logging?
3. Documentation (10%):
 - Are the setup instructions clear?
 - Can the reviewer easily run the code locally?
4. Bonus Tasks (Optional) (15%):

- Logging of requests and responses.
 - Handling API failures gracefully (e.g., retry logic or informative error messages).
5. Creativity and Efficiency (10%):
- Did the candidate go beyond the basic requirements?
 - Is the solution efficient and scalable?

Deliverables:

- A zip file containing the project source code.
- Clear documentation in the README.md file on how to set up and run the application.
- Any additional notes on potential improvements or challenges faced during development.

Be prepared to answer questions about building custom GPT models, API design, and error handling strategies.